

SAI VARUN THABETI

700741122

ASSIGNMENT 7

NEURAL NETWORKS AND DEEP LEARNING

Link for the recording:

https://drive.google.com/file/d/18kIYGJuSqHEij35DITnXYEnYH9qj9VDn/view?usp=drive_link

```
Import libraries
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow.keras.datasets import mnist

from tensorflow.keras.optimizers import RMSprop
from keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalization

%matplotlib inline

1]

Extract data and train and test dataset

# cifar100 = tf.keras.datasets.cifar100
(X_train, Y_train), (X_test, Y_test) = cifar100.load_data()

2]


classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

3]

Let's look into the dataset images

plt.figure(figsize = (16,16))
for i in range(100):
    plt.subplot(10,10,1+i)
    plt.axis('off')
    plt.imshow(X_train[i], cmap = 'gray')

4]
```



Training , Validating and Splitting trained and tested data

```

15] from sklearn.model_selection import train_test_split
    x_train, x_val, y_train, y_val = train_test_split(X_train,Y_train,test_size=0.2)

16] from keras.utils.np_utils import to_categorical
    y_train = to_categorical(y_train, num_classes = 10)
    y_val = to_categorical(y_val, num_classes = 10)

17] print(x_train.shape)
    print(y_train.shape)
    print(x_val.shape)
    print(y_val.shape)
    print(X_test.shape)
    print(Y_test.shape)

.. (40000, 32, 32, 3)
   (40000, 10)
   (10000, 32, 32, 3)
   (10000, 10)
   (10000, 32, 32, 3)
   (10000, 1)

```

```

train_datagen = ImageDataGenerator(
    preprocessing_function = tf.keras.applications.vgg19.preprocess_input,
    rotation_range=10,
    zoom_range = 0.1,
    width_shift_range = 0.1,
    height_shift_range = 0.1,
    shear_range = 0.1,
    horizontal_flip = True
)
train_datagen.fit(x_train)

val_datagen = ImageDataGenerator(preprocessing_function = tf.keras.applications.vgg19.preprocess_input)
val_datagen.fit(x_val)

```

[18]

```

from keras.callbacks import ReduceLROnPlateau
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
                                             patience=3,
                                             verbose=1,
                                             factor=0.5,
                                             min_lr=0.00001)

```

[19]

We have used only 16 layers out of 19 layers in the CNN

```

vgg_model = tf.keras.applications.VGG19(
    include_top=False,
    weights=None,
    input_shape=(32,32,3),
)

vgg_model.summary()

```

[20]

```

... Output exceeds the size limit. Open the full output data in a text editor
Model: "vgg19"

```

| Layer (type) | Output Shape | Param # |
|----------------------------|---------------------|---------|
| input_1 (InputLayer) | [(None, 32, 32, 3)] | 0 |
| block1_conv1 (Conv2D) | (None, 32, 32, 64) | 1792 |
| block1_conv2 (Conv2D) | (None, 32, 32, 64) | 36928 |
| block1_pool (MaxPooling2D) | (None, 16, 16, 64) | 0 |
| block2_conv1 (Conv2D) | (None, 16, 16, 128) | 73856 |
| block2_conv2 (Conv2D) | (None, 16, 16, 128) | 147584 |
| block2_pool (MaxPooling2D) | (None, 8, 8, 128) | 0 |
| block3_conv1 (Conv2D) | (None, 8, 8, 256) | 295168 |
| block3_conv2 (Conv2D) | (None, 8, 8, 256) | 590080 |
| block3_conv3 (Conv2D) | (None, 8, 8, 256) | 590080 |
| block3_conv4 (Conv2D) | (None, 8, 8, 256) | 590080 |

```

...
Total params: 20,024,384
Trainable params: 20,024,384
Non-trainable params: 0

```

```

model = tf.keras.Sequential()
model.add(vgg_model)
model.add(Flatten())
model.add(Dense(1024, activation = 'relu'))
model.add(BatchNormalization())
model.add(Dense(1024, activation = 'relu'))
model.add(BatchNormalization())
model.add(Dense(256, activation = 'relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(10, activation = 'softmax'))

model.summary()

```

[24]

... Output exceeds the [size limit](#). Open the full output data [in a text editor](#)
Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---|-------------------|----------|
| vgg19 (Functional) | (None, 1, 1, 512) | 20024384 |
| flatten_1 (Flatten) | (None, 512) | 0 |
| dense_4 (Dense) | (None, 1024) | 525312 |
| batch_normalization_3 (Batch Normalization) | (None, 1024) | 4096 |
| dense_5 (Dense) | (None, 1024) | 1049600 |
| batch_normalization_4 (Batch Normalization) | (None, 1024) | 4096 |
| dense_6 (Dense) | (None, 256) | 262400 |
| batch_normalization_5 (Batch Normalization) | (None, 256) | 1024 |
| dropout_1 (Dropout) | (None, 256) | 0 |

...
Total params: 21,873,482
Trainable params: 21,868,874
Non-trainable params: 4,608

```
optimizer = tf.keras.optimizers.SGD(learning_rate = 0.001, momentum = 0.9)
model.compile(optimizer= optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
history = model.fit(
    train_datagen.flow(x_train, y_train, batch_size = 128),
    validation_data = val_datagen.flow(x_val, y_val, batch_size = 128),
    epochs = 100,
    verbose = 1,
    callbacks = [learning_rate_reduction]
)
```


Output exceeds the [size limit](#). Open the full output data [in a text editor](#)

```
Epoch 1/100
313/313 [=====] - 33s 106ms/step - loss: 1.0459 - accuracy: 0.6251 - val_loss: 1.1096 - val_accuracy: 0.6193 - lr: 1.0000e-05
Epoch 2/100
313/313 [=====] - 32s 102ms/step - loss: 1.0524 - accuracy: 0.6243 - val_loss: 1.1004 - val_accuracy: 0.6269 - lr: 1.0000e-05
Epoch 3/100
313/313 [=====] - 32s 101ms/step - loss: 1.0464 - accuracy: 0.6241 - val_loss: 1.0950 - val_accuracy: 0.6269 - lr: 1.0000e-05
Epoch 4/100
313/313 [=====] - 33s 105ms/step - loss: 1.0410 - accuracy: 0.6277 - val_loss: 1.1344 - val_accuracy: 0.6244 - lr: 1.0000e-05
Epoch 5/100
313/313 [=====] - 32s 102ms/step - loss: 1.0508 - accuracy: 0.6237 - val_loss: 1.0986 - val_accuracy: 0.6256 - lr: 1.0000e-05
Epoch 6/100
313/313 [=====] - 33s 104ms/step - loss: 1.0493 - accuracy: 0.6246 - val_loss: 1.0882 - val_accuracy: 0.6293 - lr: 1.0000e-05
Epoch 7/100
313/313 [=====] - 32s 101ms/step - loss: 1.0469 - accuracy: 0.6259 - val_loss: 1.0865 - val_accuracy: 0.6267 - lr: 1.0000e-05
Epoch 8/100
313/313 [=====] - 33s 104ms/step - loss: 1.0476 - accuracy: 0.6242 - val_loss: 1.0979 - val_accuracy: 0.6231 - lr: 1.0000e-05
Epoch 9/100
313/313 [=====] - 33s 105ms/step - loss: 1.0445 - accuracy: 0.6250 - val_loss: 1.1181 - val_accuracy: 0.6211 - lr: 1.0000e-05
Epoch 10/100
313/313 [=====] - 33s 105ms/step - loss: 1.0398 - accuracy: 0.6288 - val_loss: 1.1228 - val_accuracy: 0.6246 - lr: 1.0000e-05
Epoch 11/100
313/313 [=====] - 33s 105ms/step - loss: 1.0376 - accuracy: 0.6305 - val_loss: 1.1214 - val_accuracy: 0.6233 - lr: 1.0000e-05
Epoch 12/100
313/313 [=====] - 33s 105ms/step - loss: 1.0450 - accuracy: 0.6271 - val_loss: 1.0895 - val_accuracy: 0.6288 - lr: 1.0000e-05
Epoch 13/100
...
Epoch 99/100
313/313 [=====] - 33s 105ms/step - loss: 1.0009 - accuracy: 0.6400 - val_loss: 1.0980 - val_accuracy: 0.6304 - lr: 1.0000e-05
Epoch 100/100
313/313 [=====] - 33s 104ms/step - loss: 0.9992 - accuracy: 0.6413 - val_loss: 1.1563 - val_accuracy: 0.6038 - lr: 1.0000e-05
```

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.figure()
plt.plot(acc,color = 'purple',label = 'Training Accuracy')
plt.plot(val_acc,color = 'blue',label = 'Validation Accuracy')
plt.legend()
```

<matplotlib.legend.Legend at 0x7f75101e8160>



```

loss = history.history['loss']
val_loss = history.history['val_loss']

plt.figure()
plt.plot(loss,color = 'green',label = 'Training Loss')
plt.plot(val_loss,color = 'red',label = 'Validation Loss')
plt.legend()

```

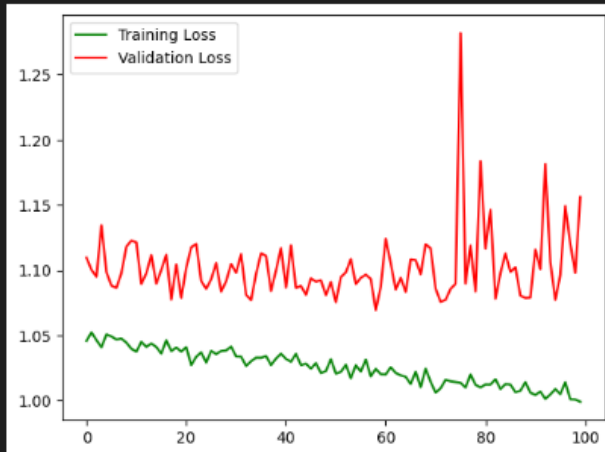
[77]

```

... <matplotlib.legend.Legend at 0x7f75101e8d30>

```

</>



```

X_test = tf.keras.applications.vgg19.preprocess_input(X_test)
y_pred = np.argmax(model.predict(X_test), axis=-1)

y_pred[:10]

```

[78]

```

... 313/313 [=====] - 3s 9ms/step

```

```

array([5, 1, 5, 5, 5, 5, 7, 5, 5, 7])

```

```

from sklearn.metrics import confusion_matrix, accuracy_score
print('Testing Accuracy : ', accuracy_score(Y_test, y_pred))

```

[79]

```

... Testing Accuracy : 0.1326

```

```
cm = confusion_matrix(Y_test, y_pred)
cm
```

[80]

```
... array([[ 7, 36,  0,  4,  0, 433,  0, 484,  0, 36],
 [ 61, 141,  0,  2,  1, 250,  0, 399,  0, 146],
 [  1,  0,  0, 10,  0, 737,  0, 250,  0,  2],
 [  0,  1,  0,  8,  3, 685,  0, 295,  0,  8],
 [  0,  0,  0, 16,  3, 779,  0, 197,  0,  5],
 [  1,  0,  0, 19,  2, 684,  0, 290,  0,  4],
 [  3,  5,  0,  9,  1, 716,  0, 252,  0, 14],
 [  1,  5,  0, 18,  2, 597,  0, 334,  0, 43],
 [ 15, 34,  0,  3,  0, 469,  0, 423,  0, 56],
 [ 27, 149,  0,  2,  1, 319,  0, 353,  0, 149]])
```

```
import itertools
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Greens):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=30)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype("float") / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    #print(cm)

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```


[81]

```
plt.figure(figsize=(8,8))
plot_confusion_matrix(cm, classes)
```

[82]

```
''' Confusion matrix, without normalization
```

```
</>
```




```
import keras
from keras.models import Sequential
from keras.preprocessing import image
from keras.layers import Activation,Dense,Dropout,Conv2D,Flatten,MaxPooling2D,BatchNormalization
from keras.datasets import cifar10
from keras import optimizers
from matplotlib import pyplot as plt
```

[8]

```
#generate cifar10 data
(x_train,y_train),(x_test,y_test) = cifar10.load_data()
```

[10]

```
#config parameters
num_classes = 10
input_shape = x_train.shape[1:4]
optimizer = optimizers.Adam(lr=0.001)
```

[9]

```
#convert label to one-hot
one_hot_y_train = keras.utils.to_categorical(y_train,num_classes=num_classes)
one_hot_y_test = keras.utils.to_categorical(y_test,num_classes=num_classes)
```


[11]

```
# check data
plt.imshow(x_train[1])
print(x_train[1].shape)
```

[12]

... (32, 32, 3)

</>



```

# build model(similar to VGG16, only change the input and output shape)
model = Sequential()
model.add(Conv2D(64,(3,3),activation='relu',input_shape=input_shape,padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(4096,activation='relu'))
model.add(Dense(2048, activation='relu'))
model.add(Dense(1024, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

```

13]

```

model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])

```

14]

```
model.summary()
```

[15]

... Output exceeds the [size limit](#). Open the full output data [in a text editor](#)
Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---|---------------------|---------|
| ===== | | |
| conv2d_1 (Conv2D) | (None, 32, 32, 64) | 1792 |
| batch_normalization (Batch Normalization) | (None, 32, 32, 64) | 256 |
| conv2d_2 (Conv2D) | (None, 32, 32, 64) | 36928 |
| batch_normalization_1 (Batch Normalization) | (None, 32, 32, 64) | 256 |
| max_pooling2d (MaxPooling2D) | (None, 16, 16, 64) | 0 |
| dropout (Dropout) | (None, 16, 16, 64) | 0 |
| conv2d_3 (Conv2D) | (None, 16, 16, 128) | 73856 |
| batch_normalization_2 (Batch Normalization) | (None, 16, 16, 128) | 512 |
| conv2d_4 (Conv2D) | (None, 16, 16, 128) | 147584 |
| ... | | |
| Total params: 27,331,914 | | |
| Trainable params: 27,323,466 | | |
| Non-trainable params: 8,448 | | |

```
history = model.fit(x=x_train, y=one_hot_y_train, batch_size=128, epochs=30, validation_split=0.1)
```

... Output exceeds the [size limit](#). Open the full output data [in a text editor](#)

```
Epoch 1/30
352/352 [=====] - 29s 81ms/step - loss: 1.7083 - accuracy: 0.3346 - val_loss: 2.6400 - val_accuracy: 0.3686
Epoch 2/30
352/352 [=====] - 27s 76ms/step - loss: 1.3159 - accuracy: 0.5254 - val_loss: 1.6935 - val_accuracy: 0.5590
Epoch 3/30
352/352 [=====] - 27s 77ms/step - loss: 1.0110 - accuracy: 0.6538 - val_loss: 1.0628 - val_accuracy: 0.6556
Epoch 4/30
352/352 [=====] - 27s 78ms/step - loss: 0.8493 - accuracy: 0.7144 - val_loss: 0.9995 - val_accuracy: 0.6856
Epoch 5/30
352/352 [=====] - 28s 79ms/step - loss: 0.7343 - accuracy: 0.7566 - val_loss: 0.8316 - val_accuracy: 0.7326
Epoch 6/30
352/352 [=====] - 28s 79ms/step - loss: 0.6515 - accuracy: 0.7907 - val_loss: 0.8145 - val_accuracy: 0.7648
Epoch 7/30
352/352 [=====] - 28s 80ms/step - loss: 0.5708 - accuracy: 0.8168 - val_loss: 0.7658 - val_accuracy: 0.7716
Epoch 8/30
352/352 [=====] - 28s 79ms/step - loss: 0.5135 - accuracy: 0.8339 - val_loss: 0.6754 - val_accuracy: 0.7944
Epoch 9/30
352/352 [=====] - 28s 79ms/step - loss: 0.4665 - accuracy: 0.8506 - val_loss: 0.7615 - val_accuracy: 0.7686
Epoch 10/30
352/352 [=====] - 28s 80ms/step - loss: 0.4213 - accuracy: 0.8662 - val_loss: 0.7079 - val_accuracy: 0.8050
Epoch 11/30
352/352 [=====] - 28s 79ms/step - loss: 0.3707 - accuracy: 0.8814 - val_loss: 0.6174 - val_accuracy: 0.8214
Epoch 12/30
352/352 [=====] - 28s 80ms/step - loss: 0.3220 - accuracy: 0.8959 - val_loss: 0.6657 - val_accuracy: 0.8202
Epoch 13/30
...
Epoch 29/30
352/352 [=====] - 28s 79ms/step - loss: 0.0811 - accuracy: 0.9749 - val_loss: 0.6592 - val_accuracy: 0.8510
Epoch 30/30
352/352 [=====] - 28s 78ms/step - loss: 0.0753 - accuracy: 0.9770 - val_loss: 0.6631 - val_accuracy: 0.8642
```

+ Code

```
# evaluate
print(model.metrics_names)
model.evaluate(x=x_test, y=one_hot_y_test, batch_size=512)
```

... ['loss', 'accuracy']
20/20 [=====] - 5s 132ms/step - loss: 0.6627 - accuracy: 0.8592

```
[0.6626988053321838, 0.8592000007629395]
```

```

model.save("keras-VGG16-cifar10.h5")
plt.imshow(x_test[1000])


result = model.predict(x_test[1000:1001]).tolist()
predict = 0
expect = y_test[1000][0]
for i, _ in enumerate(result[0]):
    if result[0][i] > result[0][predict]:
        predict = i
print("predict class:",predict)
print("expected class:",expect)

```

```

... 1/1 [=====] - 1s 740ms/step
predict class: 5
expected class: 5

```



```

# save model
model.save("keras-VGG16-cifar10.h5")


```

[16]

```

#plot the training and validation accuracy
plt.plot(history.history['accuracy'], label='training accuracy')
plt.plot(history.history['val_accuracy'], label='validation accuracy')
plt.title('Training and validation accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

```



```
#plot the training and validation loss
plt.plot(history.history['loss'], label='training loss')
plt.plot(history.history['val_loss'], label='validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```




```
import numpy as np
from sklearn.metrics import confusion_matrix

# calculate the confusion matrix
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = y_test.ravel()
cm = confusion_matrix(y_true, y_pred_classes)

# plot the confusion matrix
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion matrix')
plt.colorbar()
tick_marks = np.arange(num_classes)
plt.xticks(tick_marks, range(num_classes))
plt.yticks(tick_marks, range(num_classes))
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.show()

# plot a histogram of the predicted probabilities for a sample image
plt.hist(y_pred[1000])
plt.title('Predicted probabilities')
plt.xlabel('Probability')
plt.ylabel('Frequency')
plt.show()
```

</>



</>

